

ADVERSARIAL SEARCH

Udacity Project #3

Kyle Topasna

Problem:

This study looks at the development of a heuristic for an AI agent in the game of Knights Isolation on a 9x11 grid. The rules and description for the game can be found at here:

<https://github.com/alekzandr/ai-nanodegree/tree/master/Adversarial%20Search>

Baseline Heuristic:

The baseline heuristic uses the number of liberties available to the user's agents minus the number of liberties available to the opponent. Below is a Pseudocode representation for the baseline heuristic.

```
Define count_moves(game_state, player):
    player_location <- game_state.get_player_location(player)
    return length(game_state.get_liberties(player_location))

Define baseline_heuristic(game_state):
    player_0_moves <- count_moves(game_state, 0)
    player_1_moves <- count_moves(game_state, 1)
    return player_0_moves - player_1_moves
```

Custom Heuristics:

For the study we use two variations of the Manhattan distance for our heuristic, Max Distance and Min Distance. The Max Distance variation will favor actions that place the agent farther away from the opponent. Conversely, the Min Distance variation will favor actions that place the agent closer to the opponent.

```
Define max_distance(game_state):
    player_loc <- game_state.get_location(player)
    opp_loc <- game_state.get_location(opponent)
    return  $\sqrt{(\text{player\_loc}[x] - \text{opp\_loc}[x])^2 + (\text{player\_loc}[y] - \text{opp\_loc}[y])^2}$ 

Define min_distance(game_state):
    player_loc <- game_state.get_location(player)
    opp_loc <- game_state.get_location(opponent)
    return  $-\sqrt{(\text{player\_loc}[x] - \text{opp\_loc}[x])^2 + (\text{player\_loc}[y] - \text{opp\_loc}[y])^2}$ 
```

Baseline Results:

Here we collect results from 2 rounds of 50 games using the baseline heuristic in different match parameters. Each round, the agent trades first play initiative against the opponent.

Baseline Heuristic				
Search Depth	Move Timeouts (milliseconds)	Fair Match	Number Processes	Win Percentage
1	150	Y	1	18%
2	1000	Y	1	22%
3	100000	Y	10	34.5%
4	1000000	Y	10	31%

Max Distance Heuristic Results:

Here we collect results from 2 rounds of 50 games using the custom heuristic and matching match parameters to the baseline:

Max Distance Heuristic				
Search Depth	Move Timeouts (milliseconds)	Fair Match	Number Processes	Win Percentage
1	150	Y	1	18%
2	500	Y	1	12.5%
3	10000	Y	4	23%
4	100000	Y	10	22.5%

Min Distance Heuristic Results:

Here we collect results from 2 rounds of 50 games using the custom heuristic and matching match parameters to the baseline:

Min Distance Heuristic				
Search Depth	Move Timeouts (milliseconds)	Fair Match	Number Processes	Win Percentage
1	150	Y	1	15.5%
2	500	Y	1	17%
3	10000	Y	4	39%
4	100000	Y	10	44%

We see more completeness at the expense of computational performance when combining the baseline heuristic with the developed Max Distance Heuristic as a weighted function called the aggressive greedy heuristic.

```

Define aggressive_greedy(game_state, weight_1, weight_2):
    Return (weight_1 * min_distance) + (weight_2 * baseline_heuristic)

```

Aggressive Greedy Heuristic				
Search Depth	Move Timeouts (milliseconds)	Fair Match	Number Processes	Win Percentage
1	150	Y	1	17%
2	1000	Y	1	23.5%
3	100000	Y	10	33%
4	1000000	Y	10	41.5%

Below we look at using the aggressive greedy in combination with the iterative deepening version of the minimax search. Our agent will have a max search depth of ten.

Aggressive Greedy Heuristic with Iterative Deepening				
Search Depth	Move Timeouts (milliseconds)	Fair Match	Number Processes	Win Percentage
10	150	Y	1	74%
10	1000	Y	1	79%

Project Questions:

- What features of the game does your heuristic incorporate, and why do you think those features matter in evaluating states during search?

The Max Distance, Min Distance, and Aggressive Greedy heuristics all use the distance between players as a feature for evaluating states. The Aggressive Greedy heuristic ties in a known good heuristic of counting the liberties available to each player.

- Analyze the search depth your agent achieves using your custom heuristic. Does search speed matter more or less than accuracy to the performance of your heuristic?

We see greater confidence in performance when using deeper searches. This comes at the cost of greater search time. The heuristics developed for this project rely more on accuracy than search speed. However, when adding iterative deepening to the aggressive greedy heuristic, we achieve near optimal performance at the first two levels.

Future Improvements:

In future iterations, we could implement a different search algorithm such as Monte Carlo Tree Search. Additionally, we could implement machine learning systems to learn the best plays in situations from past experience. Reinforcement Learning techniques have proved successful in these domains.